

Creating and managing cloud resources on AWS a) tour of aws

Tour of AWS

1. AWS Management Console
 - Web-based interface for managing AWS services.
 - Provides a unified dashboard for monitoring and configuration.
2. AWS Global Infrastructure
 - Regions: Physical locations with multiple Availability Zones.
 - Availability Zones (AZs): Data centers within a region for high availability.
 - Edge Locations: Used for content delivery via AWS CloudFront.
3. Core AWS Services
 - Compute: EC2, Lambda, ECS, EKS
 - Storage: S3, EBS, Glacier
 - Networking: VPC, Route 53, CloudFront
 - Databases: RDS, DynamoDB, Aurora
 - Security & IAM: IAM, KMS, WAF
4. Billing & Pricing
 - Pay-as-you-go model
 - Free Tier options available
 - Cost Explorer for expense tracking
5. Monitoring & Management
 - AWS CloudWatch: Logs and metrics monitoring
 - AWS CloudTrail: Tracks API calls for security auditing
 - AWS Config: Ensures compliance
6. Deployment & Automation
 - AWS CloudFormation: Infrastructure as Code (IaC)
 - AWS Elastic Beanstalk: Automated deployment
 - AWS Lambda: Serverless computing
7. Security Best Practices
 - Use IAM roles and policies
 - Enable MFA for accounts
 - Encrypt data in transit and at rest

IAM

AWS Identity and Access

Management (IAM) – Steps

1. Access AWS IAM
 - Sign in to AWS Management Console.
 - Navigate to IAM service.
2. Create a New User
 - Click Users → Add user.
 - Enter Username (e.g., test-user).
 - Select Access type:
 - Programmatic Access (for CLI/SDK/API).
 - AWS Management Console (set a custom password).
 - Click Next: Permissions.
3. Assign Permissions
 - Add user to Group (assign policies like AdministratorAccess, AmazonS3FullAccess).
 - OR attach policies directly to the user.
 - Click Next: Tags.
4. Add Tags (Optional)
 - Add key-value pairs for organization.
 - Click Next: Review.
5. Review and Create User
 - Verify details, click Create user.
 - Download .csv file (contains Access Key ID & Secret Access Key).
6. Create and Attach a Role
 - Go to Roles → Click Create role.
 - Select AWS Service (e.g., EC2).
 - Attach necessary policies (e.g., AmazonS3FullAccess).
 - Name the role and create it.
 - Attach the role to an AWS service (e.g., EC2 instance).
7. Review IAM Policies and Groups
 - Regularly review permissions to ensure least privilege access.

INTRODUCTION TO AMAZON EC2

Getting Started with Amazon EC2 (Elastic Compute Cloud)

Amazon EC2 is a web service that provides resizable compute capacity in the cloud, enabling developers to scale their applications easily.

Steps to Get Started with Amazon EC2:

1. Access the AWS Management Console

- Sign in to the AWS Management Console.
- Navigate to EC2 under the "Compute" section.

2. Launch an EC2 Instance

- Click on **Launch Instance**.
- **Choose an Amazon Machine Image (AMI):**
 - Select an AMI from the list (includes OS and pre-installed applications).
 - Popular choices: **Amazon Linux 2, Ubuntu, Windows Server**.
- **Choose an Instance Type:**
 - Select an instance type based on your needs.
 - Common types: **t2.micro (free-tier eligible), m5.large**, etc.
 - Click **Next: Configure Instance Details**.
- **Configure Instance Details:**
 - Set the number of instances, network settings, IAM roles, and other advanced options.
 - Click **Next: Add Storage**.
- **Add Storage:**
 - Define the size and type of storage (EBS volumes) to attach.
 - Default settings are sufficient for basic setups.
 - Click **Next: Add Tags**.
- **Add Tags:**
 - Optionally, add tags (key-value pairs) to help manage and identify your instance.
 - Click **Next: Configure Security Group**.
- **Configure Security Group:**
 - A security group acts as a virtual firewall.
 - Set rules to control traffic to your instance.
 - For basic setups, add an SSH rule to allow access from your IP.
 - Click **Review and Launch**.
- **Review and Launch:**
 - Review instance settings and click **Launch**.
 - Create or select an SSH key pair for secure access.
 - Choose **Create a new key pair** or **Use an existing one**, then click **Launch Instances**.

3. Connect to Your EC2 Instance

- Once the instance is running, click **View Instances** to see details.

INTRODUC TO AMAZON SIMPLE STORAGE SERVICE S3

Amazon S3: Getting Started Guide

1. Access the AWS Management Console

- Sign in to the AWS Management Console.
- Navigate to S3 under the "Storage" section.

2. Create an S3 Bucket

- Click on Create bucket.
- Bucket Name: Enter a globally unique name for your bucket.
- Region: Choose an AWS Region where the bucket will reside.
- Bucket Settings: Configure settings like versioning, server access logging, and encryption (these can be adjusted later).
- Block Public Access: By default, S3 buckets block public access. Adjust if needed.
- Click Create Bucket.

3. Upload Objects to S3

- Click on your newly created bucket.
- Click Upload.
- Drag and drop files or use the file picker to select files from your local machine.
- Set object properties like storage class and encryption if needed.
- Click Upload to store the files in S3.

4. Manage Bucket Permissions

- Go to the Permissions tab in your bucket.
- Manage access using:
 - Bucket policies
 - Access control lists (ACLs)
 - IAM roles
- To make files public, modify the bucket policy or set specific permissions on individual objects.

5. Enable Versioning

- Versioning helps keep multiple versions of an object in one bucket. To enable:
 - Go to Bucket Settings → Properties → Versioning.
 - Click Edit and enable versioning.

6. Set Up Lifecycle Policies

- Lifecycle policies allow automatic transitioning of objects between storage classes or deletion after a period.
- Go to the Management tab and click Create lifecycle rule.
- Define rules to transition objects to cheaper storage classes (e.g., Glacier) or delete them after a set time.

7. Access Objects in S3

- Objects can be accessed via the AWS S3 Console or programmatically using the AWS SDKs or CLI.
- To generate a pre-signed URL for temporary access:
- `aws s3 presign s3://bucket-name/object-key --expires-in 3600`

8. Delete S3 Bucket

- To delete a bucket, first delete all objects within it.
- Then, delete the bucket itself through the AWS Console or CLI.

CREATING AND RUNNING CONTAINERS

A Container That Can Run Consistently Across Different Environments

Steps to Create and Run Containers:

1. Install Docker:

- Install Docker by following the instructions on the Docker website.
- Verify the installation:
`docker --version`

2. Create a Dockerfile:

- A Dockerfile is a text document that contains all the commands to assemble a container image.
- Example Dockerfile for a simple Node.js application:
`# Use an official Node.js runtime as a parent image`
`FROM node:14`
`# Set the working directory in the container`
`WORKDIR /usr/src/app`
`# Copy the local files to the container`
`COPY package*.json ./`
`# Install any needed dependencies`
`RUN npm install`
`# Copy the rest of the application code`
`COPY . .`
`# Make port 8080 available to the world outside this container`
`EXPOSE 8080`
`# Run the app when the container launches`
`CMD ["node", "app.js"]`

3. Build the Docker Image:

- Navigate to the directory containing the Dockerfile and run:
`docker build -t my-node-app .`
- This command builds the Docker image and tags it as my-node-app.

4. Run the Docker Container:

- Use the following command to run the container:
`docker run -p 8080:8080 my-node-app`
- The `-p 8080:8080` flag maps port 8080 on your local machine to port 8080 in the container.

5. Manage Containers:

- To list running containers:
`docker ps`
- To stop a container:
`docker stop <container_id>`
- To remove a container:
`docker rm <container_id>`
- To remove an image:
`docker rmi <image_id>`

6. Using Docker Compose (Optional):

- For complex applications with multiple services, use

- To connect via SSH:

- Select the instance, click **Connect**.
- Follow the provided instructions, typically using:
`ssh -i /path/to/key-pair.pem ec2-user@your-ec2-public-dns`

- For Windows instances, connect using RDP (Remote Desktop Protocol).

4. Manage and Monitor the Instance

- Use **Instance State** actions to **start, stop, reboot, or terminate** the instance.
- Monitor performance via **CloudWatch** (CPU usage, network traffic, etc.).

5. Terminate the Instance

- To avoid charges, **terminate** the instance when it is no longer needed.

AMAZON REALTIONAL DATABASE SERVICE

Amazon RDS (Relational Database Service) makes it easy to set up, operate, and scale a relational database in the cloud. With RDS, you can run several database engines, including SQL Server, MySQL, PostgreSQL, Oracle, and MariaDB.

Steps to Get Started with Amazon RDS - SQL Server:

1. **Access the AWS Management Console**
 - Sign in to the AWS Management Console.
 - Navigate to RDS under the "Database" section.
2. **Create a SQL Server RDS Instance**
 - Click on **Create database**.
 - **Database Creation Method:**
 - Choose either **Standard Create** for detailed setup options or **Easy Create** for basic configurations.
 - **Engine Options:**
 - Select **Microsoft SQL Server** from the list of database engines.
 - Choose the specific SQL Server edition (e.g., Express, Web, Standard, or Enterprise).
 - **Version:**
 - Select the version of SQL Server you want to use (e.g., SQL Server 2019).
 - **Instance Specifications:**
 - Choose the **DB instance class** based on your performance needs (e.g., db.t3.micro for a free-tier eligible instance or db.m5.large for higher performance).
 - **Storage:**
 - Set the allocated storage size. You can also enable **storage autoscaling** to allow the database to grow automatically as needed.
 - **Database Settings:**
 - Set the **DB Instance Identifier** (name of your database instance).
 - Define the **Master Username** and **Master Password**.
 - **VPC and Network Settings:**

Docker Compose to manage multi-container Docker applications.

- Example docker-compose.yml file:
- version: '3'
- services:
- web:
- image: my-node-app
- ports:
- - "8080:8080"
- redis:
- image: "redis:alpine"
- Start the services with:
- docker-compose up

MANAGING RESOURCES USING TERRAFORM:

Managing Resources Across Various Cloud Providers Using Terraform

Steps to Manage Resources Using Terraform:

1. Install Terraform:

- Download and install Terraform from the official website.
- Verify the installation:
- terraform --version

2. Create a Terraform Configuration File:

- Define your infrastructure in a .tf file. For example, to create an EC2 instance in AWS:
- provider "aws" {
- region = "us-west-2"
- }
- resource "aws_instance" "example"
- {
- ami = "ami-0c55b159cbf1e1f0"
- instance_type = "t2.micro"
- }
- tags = {
- Name = "ExampleInstance"
- }
- }
- This example sets up an AWS provider and defines a single EC2 instance.

3. Initialize Terraform:

- Run the following command to initialize Terraform. This command downloads the required provider plugins:
- terraform init

4. Preview the Changes:

- Before applying the changes, you can preview what Terraform will do:
- terraform plan

5. Apply the Configuration:

- To create the resources defined in your configuration file, run:
- terraform apply
- Terraform will prompt you to confirm the action. Type yes to proceed.

6. Managing Resources:

- Terraform can be used to manage the entire lifecycle of your infrastructure, including updates and deletions.
- To update resources, modify the .tf file and run:
- terraform apply
- To delete resources, use:
- terraform destroy

<div><ul style="list-style-type: none"><ul style="list-style-type: none">■ Select the VPC, subnet, and security group settings that will govern network access to your instance.○ Backup and Maintenance:<ul style="list-style-type: none">■ Configure automated backups, backup retention period, and the preferred backup window.■ Set up a maintenance window if you want to specify a particular time for updates and patches.○ Click Create database to start the creation process.<p>3. Connect to the SQL Server RDS Instance</p><ul style="list-style-type: none">○ Once the RDS instance is available (this may take a few minutes), you can connect to it using SQL Server Management Studio (SSMS) or another SQL client.○ Endpoint:<ul style="list-style-type: none">■ Find the endpoint and port in the RDS console under the Connectivity & security tab.■ Use these details in your SQL client to connect.○ Connect via SSMS:<ul style="list-style-type: none">■ Open SSMS, enter the RDS endpoint as the Server name.■ Choose SQL Server Authentication and enter the master username and password.<p>4. Manage and Monitor the RDS Instance</p><ul style="list-style-type: none">○ Use the RDS Dashboard to monitor the instance's performance, including CPU usage, memory, and storage.○ Configure CloudWatch Alarms to alert you if the database goes beyond certain thresholds (e.g., high CPU usage or low available storage).○ Perform administrative tasks like creating snapshots (manual backups), restoring backups, or scaling the instance size.<p>5. Backup and Restore</p><ul style="list-style-type: none">○ Automated Backups: Ensure automated backups are enabled. You can configure the backup retention period.</div>		
---	--	--

<ul style="list-style-type: none">○ Manual Snapshots: Create a manual snapshot to capture the current state of your database. This snapshot can be used later for restoring or replicating the database.○ Restore a Database: To restore, go to the Snapshots tab, select the snapshot, and choose Restore Snapshot. <p>6. Security and Access Control</p> <ul style="list-style-type: none">○ Use AWS IAM to control who can manage the RDS instances.○ Ensure encryption is enabled for data at rest by using AWS-managed keys.○ Use SSL/TLS to secure data in transit between your application and the RDS instance. <p>7. Terminate the RDS Instance</p> <ul style="list-style-type: none">○ When you no longer need the RDS instance, terminate it to avoid charges.○ Go to the RDS Dashboard, select the instance, and click Delete.○ You can choose to retain or delete the final snapshot during the deletion process.		